



Intelligent Systems for Modern Spaceflight

- Dheepak karan

Executive Summary

This project conducted an in-depth analysis of SpaceX's launch data to forecast the success of first-stage landings using machine learning techniques. By utilizing historical records, we implemented logistic regression, support vector machines, and k-nearest neighbors, achieving accuracy rates exceeding 80%. Performance was further improved through hyperparameter tuning with GridSearchCV, highlighting the value of AutoML in predictive modeling. The analysis also uncovered key insights, such as the lack of correlation between payload mass and landing success, and demonstrated the consistency of SpaceX's performance across different payload sizes. These results can support strategic advancements in future launch missions.

Introduction to Commercial Space Travel

Commercial space travel represents a bold new era, driven by private companies expanding human access beyond Earth. Key players in this field include:

Virgin Galactic: Focused on suborbital space tourism.

Rocket Lab: Specializes in deploying small satellites into orbit.

Blue Origin: Advancing reusable rocket technology for both suborbital and orbital missions.

SpaceX: At the forefront with breakthroughs like ISS missions, the Starlink satellite network, and crewed commercial flights.

Spotlight on SpaceX: Renowned for its reusable Falcon 9 rockets, SpaceX has dramatically reduced launch costs—around \$62 million per mission compared to traditional rates reaching \$165 million—setting new efficiency standards in the industry.

Role of a Data Scientist in Space Y

Introducing Space Y: An emerging aerospace company with ambitions to compete with industry leaders like SpaceX.

Data Scientist's Role:

Goal: Evaluate and enhance launch cost efficiency.

Approach: Gather and examine data related to SpaceX's operations, and build interactive dashboards to visualize key cost indicators.

Innovation: Create a machine learning model to predict the likelihood of Falcon 9 first-stage reusability using publicly available datasets.

Rocket Stage Analogy for Better Understanding:

First Stage – The “Power Booster”: Like the thick, sturdy bottom of a pencil, this stage provides the massive thrust needed for liftoff.

Second Stage – The “Precision Guide”: Comparable to the pencil's mid-section, it's responsible for accurately placing the payload into its intended orbit.

Data Collection Using SpaceX API methodology

Exploring the SpaceX API:

The SpaceX REST API provides extensive information on past and present launches, including rocket specifications and mission outcomes, making it a valuable resource for data-driven analysis.

API Overview:

Base URL: <https://api.spacexdata.com/v4/>

Key Endpoints:

/capsules: Provides capsule-related details.

/cores: Offers data on rocket cores.

/launches/past: Central to this project, it supplies historical launch data for analysis.

Data Retrieval Approach:

Python's requests library is used to send GET requests to the /launches/past endpoint.

The API returns responses in JSON format, which are parsed and processed for further analysis.

Supplementary Methods:

To enhance the dataset, additional Falcon 9 information is extracted from Wikipedia using Python's BeautifulSoup library for web scraping.

Data Wrangling for SpaceX Launch Analysis Methodology

Processing JSON Data:

To enable efficient analysis, nested JSON responses from the API are flattened and converted into a pandas DataFrame using `json_normalize`.

Data Cleaning Workflow:

Handling Missing Values: Fill missing entries in the PayloadMass column using the column's mean value.

Removing Irrelevant Columns: Drop fields like LandingPad if they are not relevant to the current analysis.

Filtering Records: Exclude entries that do not involve Falcon 9 launches to maintain consistency in the dataset.

Key Features for Analysis:

Important attributes include:

Flight Number

Booster Version

Payload Mass

Launch Site

Launch Outcome

Orbit

Reused Components

Landing Type

These features are essential for assessing mission performance and identifying trends.

Outcome Classification Strategy:

Launch results are simplified into a binary classification:

1 for successful landings

0 for failures

This conversion streamlines the use of machine learning models for predictive analysis.

Exploratory Data Analysis (EDA) Methodology

What is EDA (Exploratory Data Analysis)?

EDA is a vital first phase in any data science project, focused on understanding the data's structure, patterns, and potential through statistical summaries and visual exploration.

Core EDA Practices:

- Perform basic data checks and review summary statistics.

- Use visualizations to uncover trends, anomalies, and relationships between features.

- Evaluate the data's potential to predict Falcon 9 first-stage landing outcomes using deeper statistical analysis.

Key Analytical Areas:

- Success Rate Over Time: Monitor how landing success rates evolve, using launch number as a key indicator.

- Launch Site Impact: Examine the influence of different launch sites on landing success, supported by statistical comparisons.

Preparing Data for Machine Learning:

- Select relevant features based on their correlation with the target variable (success or failure).

- Transform categorical variables using techniques like one-hot encoding to make the data suitable for machine learning algorithms.

Interactive Visual Analytics and Dashboard Building- methodology

Overview of Interactive Visual Analytics:

Interactive visual analytics involves leveraging dynamic visualization tools that let users explore and interact with data in real time, leading to deeper insights and better decision-making.

Tools and Techniques:

Folium: Ideal for creating interactive maps to visualize spatial data, such as the locations and distribution of launch sites.

Plotly Dash: A powerful Python framework for building web-based dashboards that support user interaction through components like dropdowns and sliders.

Methodology:

Part 1 – Geospatial Analysis with Folium:

Map out launch sites to explore geographic trends and assess any location-based advantages or constraints.

Part 2 – Interactive Dashboard with Plotly Dash:

Develop an interactive dashboard integrating visual components like dropdown menus, checklists, and range sliders to allow users to filter and explore launch data dynamically.

Building the Machine Learning Pipeline- methodology

Predictive Analysis Overview:

Predictive analysis involves using machine learning algorithms to forecast future outcomes based on patterns identified in historical data.

Data Preprocessing:

An essential step that prepares the dataset for effective model training by ensuring consistency and accuracy.

Key steps include:

Standardizing Numeric Features: Scale values to a consistent range for improved model performance.

Handling Missing Values: Impute or remove null entries to maintain data integrity.

Encoding Categorical Variables: Convert text-based categories into numerical format for compatibility with ML algorithms.

Data Splitting with `train_test_split`:

Training Set: Used by the model to learn from the data.

Testing Set: Held back to evaluate the model's ability to generalize and make accurate predictions on unseen data.

Model Training, Optimization, and Evaluation- methodology

Model Training Techniques:

Several machine learning algorithms were applied to classify the success of Falcon 9 first-stage landings:

Logistic Regression: Estimates the probability of binary outcomes (success or failure).

Support Vector Machines (SVM): Identifies the optimal boundary (hyperplane) that best separates the two outcome classes.

Decision Trees: Splits data based on decision rules to classify outcomes.

K-Nearest Neighbors (KNN): Predicts outcomes based on the most common class among the closest data points.

Optimizing Model Performance:

Grid Search: A systematic method to find the best combination of hyperparameters by testing multiple configurations, improving overall model accuracy.

Model Evaluation:

Confusion Matrix: A performance summary showing how well the model predicts outcomes:

True Positives (TP): Correctly predicted successes

False Positives (FP): Incorrectly predicted successes

True Negatives (TN): Correctly predicted failures

False Negatives (FN): Incorrectly predicted failures

EDA and SQL Results

Task 2

Display 5 records where launch sites begin with the string 'CCA'

```
%sql SELECT * FROM SPACEXTBL WHERE "Launch_Site" LIKE 'CCA%' LIMIT 5;
```

```
* sqlite:///my_data1.db  
one.
```

Date	Time (UTC)	Booster_Version	Launch_Site	Payload	PAYLOAD_MASS
2010-06-04	18:45:00	F9 v1.0 B0003	CCAFS LC-40	Dragon Spacecraft Qualification Unit	
2010-12-08	15:43:00	F9 v1.0 B0004	CCAFS LC-40	Dragon demo flight C1, two CubeSats, barrel of Brouere cheese	
2012-05-22	7:44:00	F9 v1.0 B0005	CCAFS LC-40	Dragon demo flight C2	
2012-10-08	0:35:00	F9 v1.0 B0006	CCAFS LC-40	SpaceX CRS-1	
2013-03-01	15:10:00	F9 v1.0 B0007	CCAFS LC-40	SpaceX CRS-2	

Task 1

Display the names of the unique launch sites in the space mission

```
%sql SELECT DISTINCT "Launch_Site" FROM SPACEXTBL;
```

```
* sqlite:///my_data1.db  
one.
```

Launch_Site
CCAFS LC-40
VAFB SLC-4E
KSC LC-39A
CCAFS SLC-40

EDA and SQL Results

Task 4

Display average payload mass carried by booster version F9 v1.1

```
%sql SELECT AVG("PAYLOAD_MASS__KG_") AS Average_Payload_Mass FROM SPACEXTBL WHERE "Booster_Version" LIKE 'F9 v1.1';
```

```
* sqlite:///my_data1.db  
one.
```

Average_Payload_Mass

2534.6666666666665

Task 3

Display the total payload mass carried by boosters launched by NASA (CRS)

```
%sql SELECT SUM("PAYLOAD_MASS__KG_") AS Total_Payload_Mass FROM SPACEXTBL WHERE "Customer" LIKE '%NASA (CRS)%';
```

```
* sqlite:///my_data1.db  
Done.
```

Total_Payload_Mass

48213

EDA and SQL Results

Task 6

List the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000

```
%sql SELECT "Booster_Version" FROM SPACEXTBL WHERE "Landing_Outcome" = 'Success (drone ship)' AND "PAYLOAD_MASS_Kg_" > 4000 AND "PAYLOAD_MASS_Kg_" < 6000;
sqlite:///my_data1.db
me.
```

Booster_Version
F9 FT B1022
F9 FT B1026
F9 FT B1021.2
F9 FT B1031.2

Task 5

List the date when the first succesful landing outcome in ground pad was acheived.

Hint: Use min function

```
%sql SELECT MIN("Date") AS First_Successful_Landing FROM SPACEXTBL WHERE "Landing_Outcome" = 'Success (ground pad)';
```

```
* sqlite:///my_data1.db
one.
```

First_Successful_Landing

2015-12-22

EDA and SQL Results

Task 7

List the total number of successful and failure mission outcomes

```
%sql SELECT CASE WHEN "Mission_Outcome" LIKE '%Success%' THEN 'Success' WHEN "Mission_Outcome" LIKE '%Failure%' THEN 'Failure' ELSE 'Other' END AS Outcome, COUNT(*) AS Total FROM
```

```
* sqlite:///my_data1.db
```

one.

Outcome	Total
---------	-------

Failure	1
---------	---

Success	100
---------	-----

Task 8

List the names of the booster_versions which have carried the maximum payload mass. Use a subquery

```
%sql SELECT "Booster_Version" FROM SPACEXTBL WHERE "PAYLOAD_MASS__KG_" = (SELECT MAX("PAYLOAD_MASS__KG_") FROM SPACEXTBL);
```

```
* sqlite:///my_data1.db  
one.
```

Booster_Version

F9 B5 B1048.4

F9 B5 B1049.4

F9 B5 B1051.3

F9 B5 B1056.4

F9 B5 B1048.5

F9 B5 B1051.4

F9 B5 B1049.5

F9 B5 B1060.2

F9 B5 B1058.3

F9 B5 B1051.6

F9 B5 B1060.3

F9 B5 B1049.7

EDA and SQL Results

Task 9

List the records which will display the month names, failure landing_outcomes in drone ship ,booster versions, launch_site for the months in year 2015.

Note: SQLite does not support monthnames. So you need to use substr(Date, 6,2) as month to get the months and substr(Date,0,5)=2015 for year.

```
%sql SELECT substr(Date, 6, 2) AS Month, "Booster_Version", "Launch_Site", "Landing_Outcome" FROM SPACEXTBL WHERE substr(Date, 1, 4) = '2015' AND "Landing_Outcome" LIKE '%Failure (dr
```

```
* sqlite:///my_data1.db
```

Done.

Month	Booster_Version	Launch_Site	Landing_Outcome
01	F9 v1.1 B1012	CCAFS LC-40	Failure (drone ship)
04	F9 v1.1 B1015	CCAFS LC-40	Failure (drone ship)

Task 10

Rank the count of landing outcomes (such as Failure (drone ship) or Success (ground pad)) between the date 2010-06-04 and 2017-03-20, in descending order.


```
%sql SELECT "Landing_Outcome", COUNT(*) AS Outcome_Count FROM SPACEXTBL WHERE "Date" BETWEEN '2010-06-04' AND '2017-03-20' GROUP BY "Landing_Outcome" ORDER BY Outcome_Count DESC;
```

```
* sqlite:///my_data1.db  
one.
```


Landing_Outcome	Outcome_Count
No attempt	10
Success (drone ship)	5
Failure (drone ship)	5
Success (ground pad)	3
Controlled (ocean)	3
Uncontrolled (ocean)	2
Failure (parachute)	2
Precluded (drone ship)	1

EDA WITH VISUALIZATION RESULTS


```
with pd.option_context('mode.use_inf_as_na', True):
```



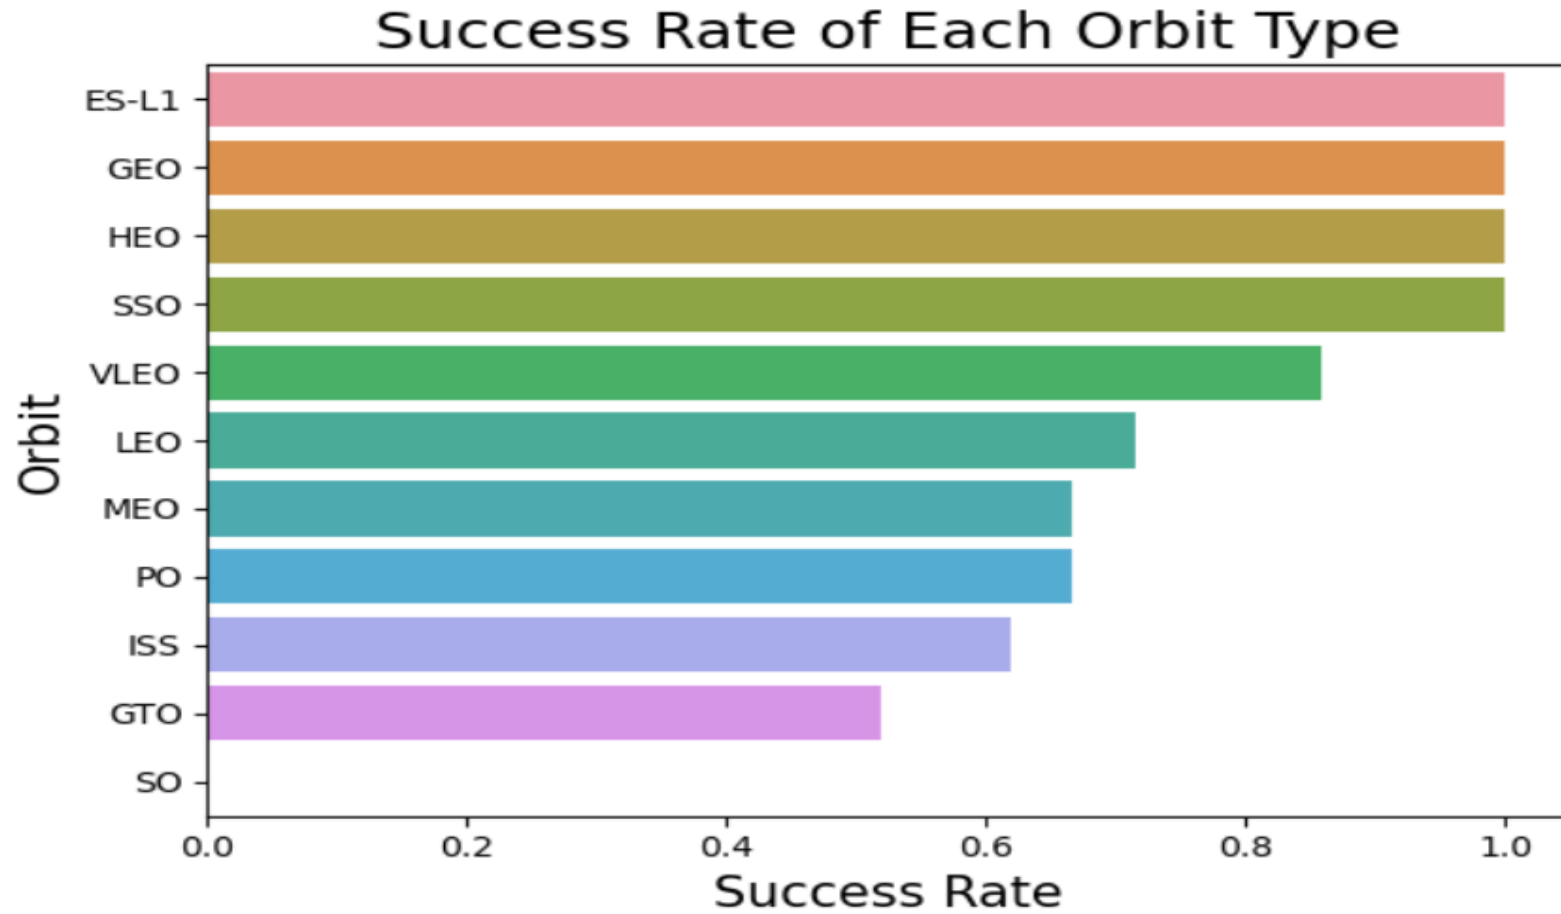
Relationship between Flight Number and Launch Site by Class




Relationship between Payload Mass and Launch Site by Class




```
sns.barplot(x=class, y=orbit, data=orbit_success_rate_sorted, ci=None)
```




Relationship between Flight Number and Orbit Type by Class



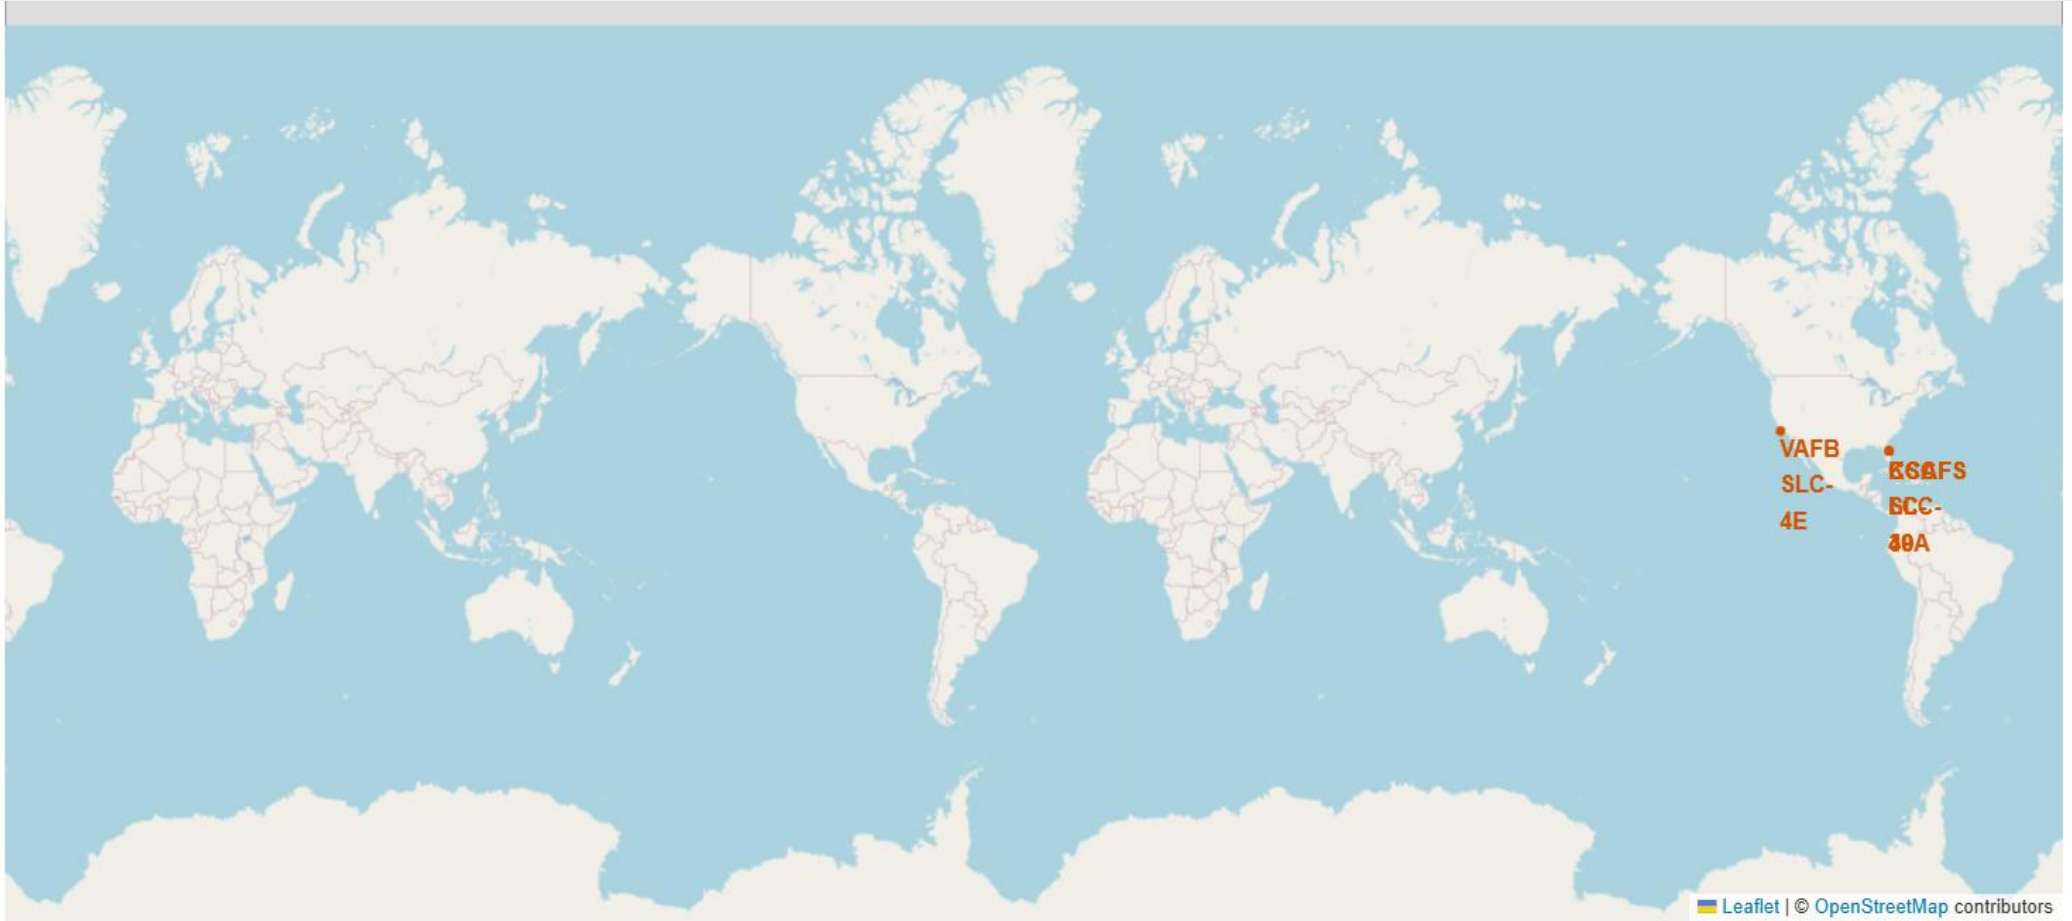
Relationship between Payload Mass and Orbit Type by Class

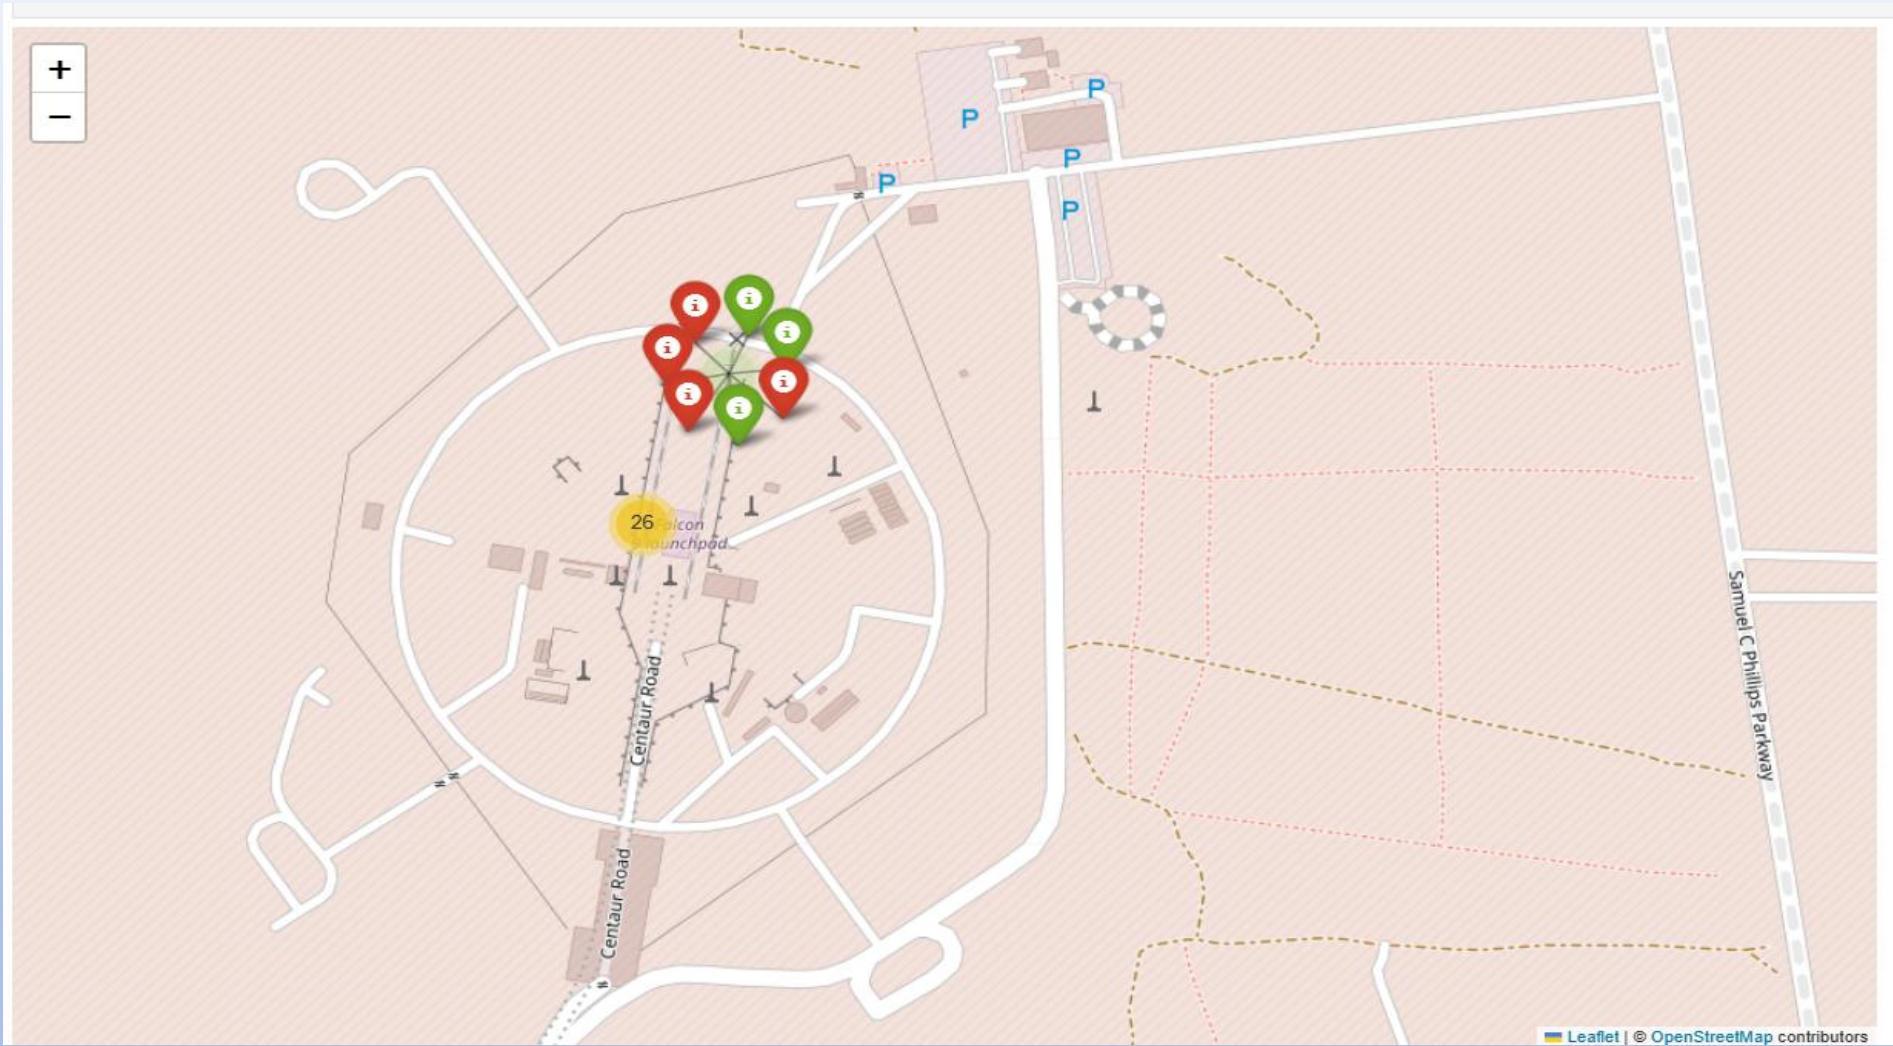


Yearly Launch Success Trend

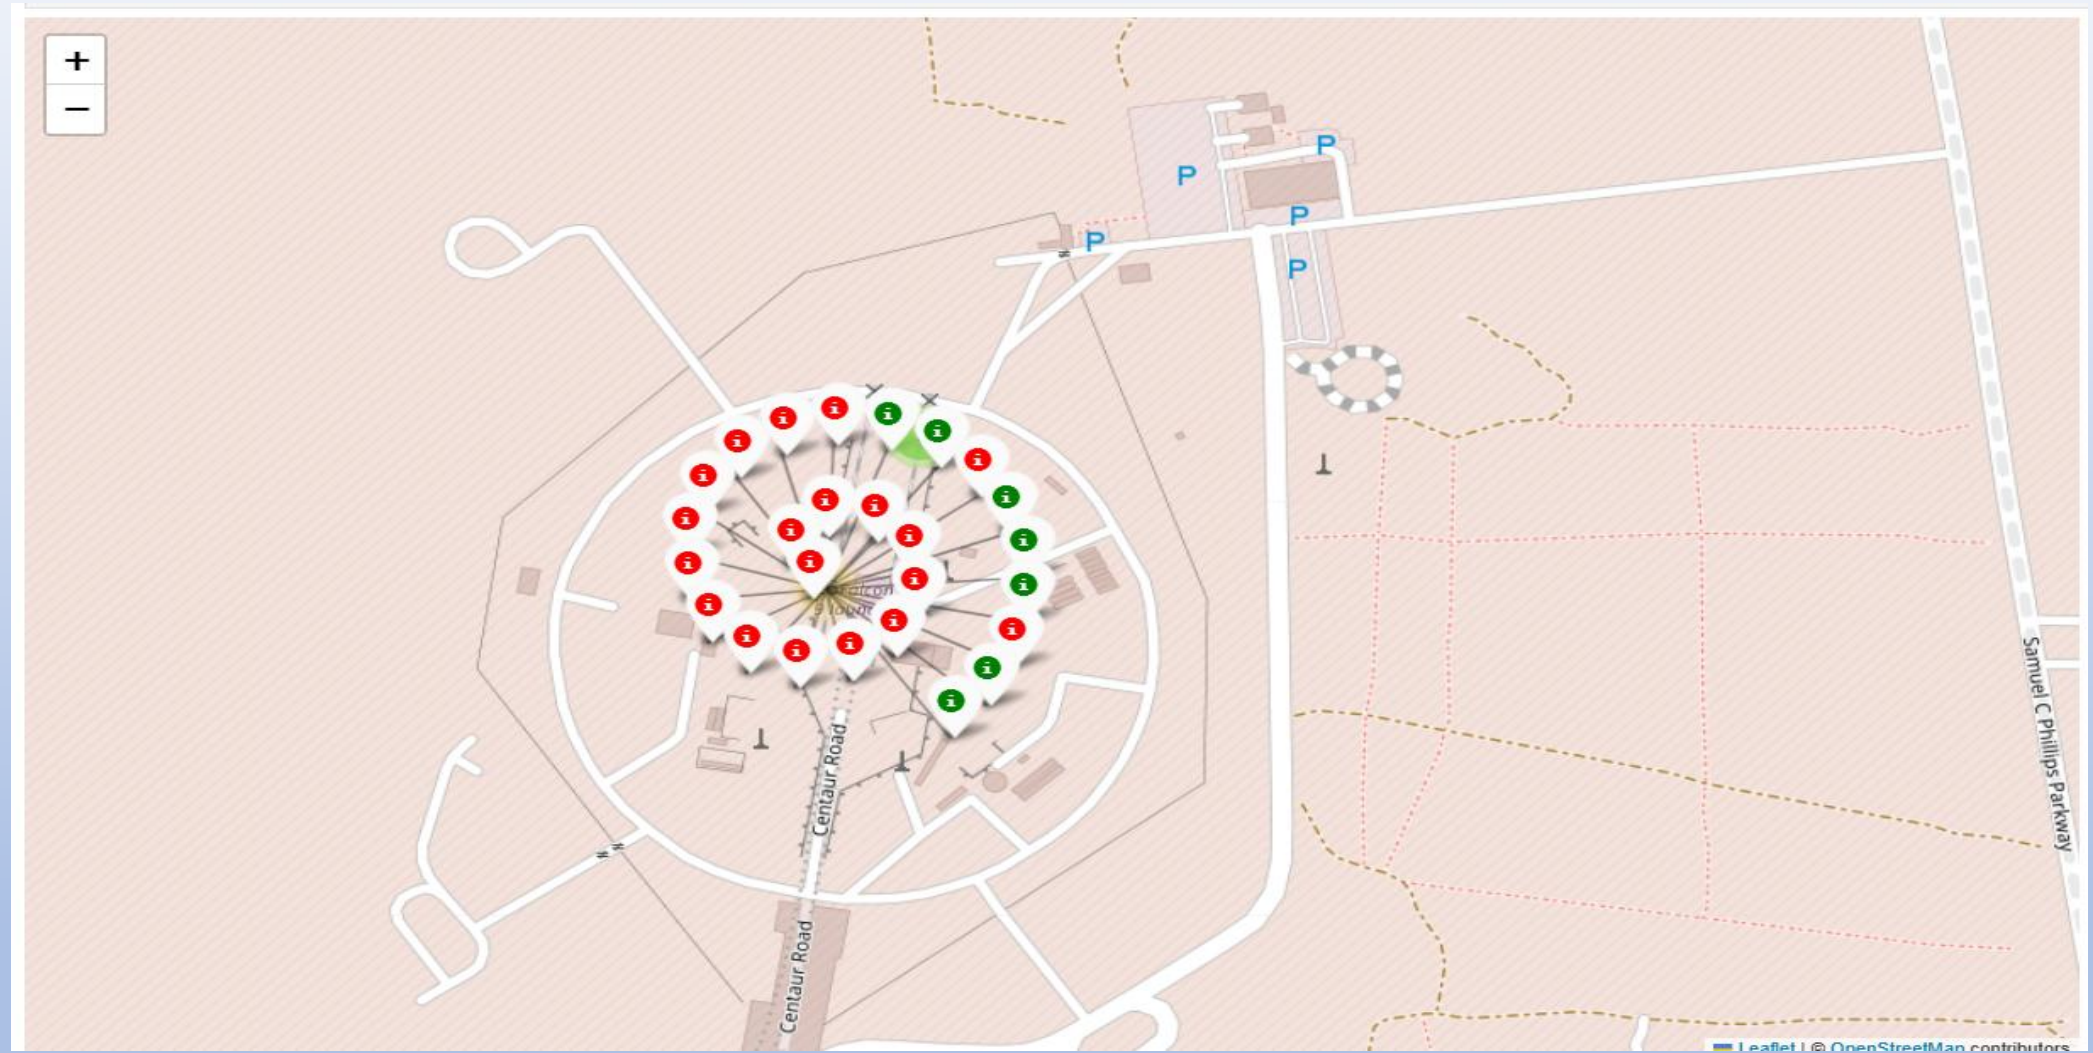


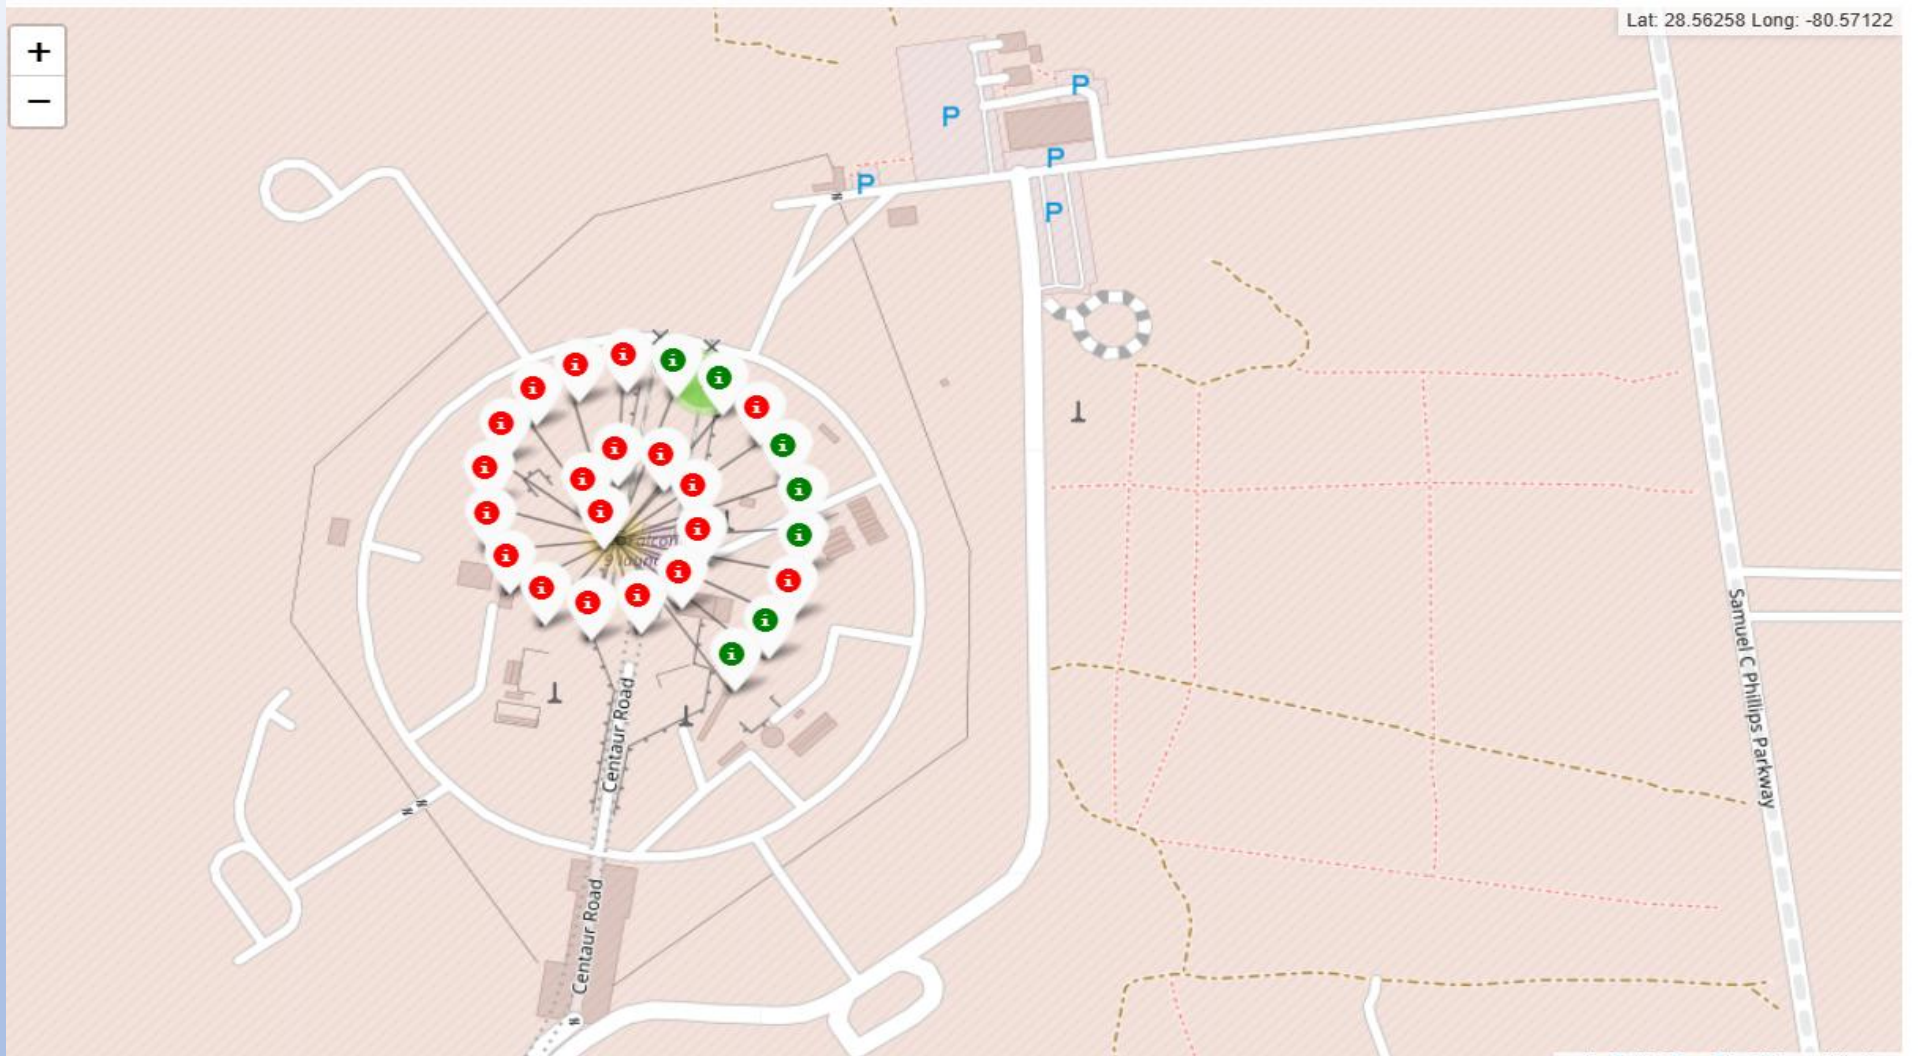
INTERACTIVE MAPS WITH FOLIUM

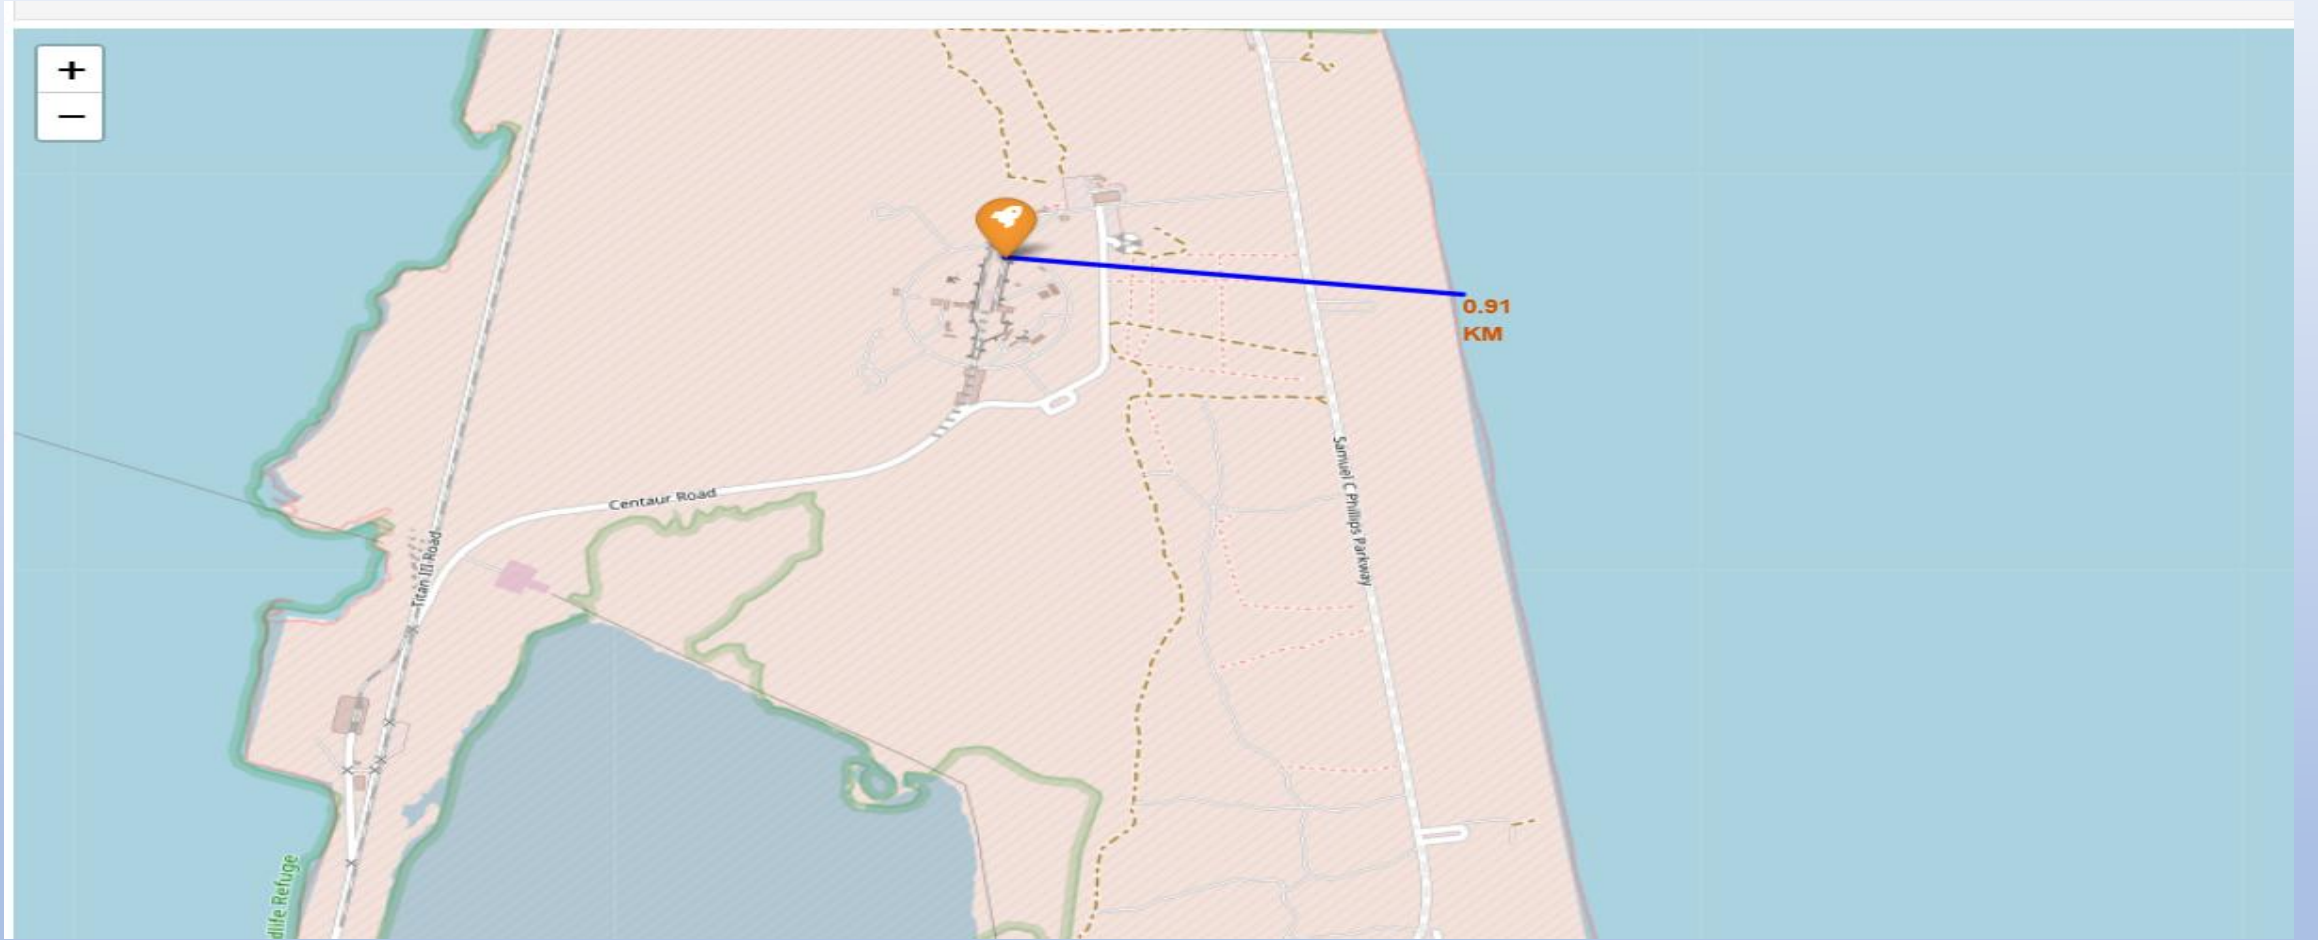


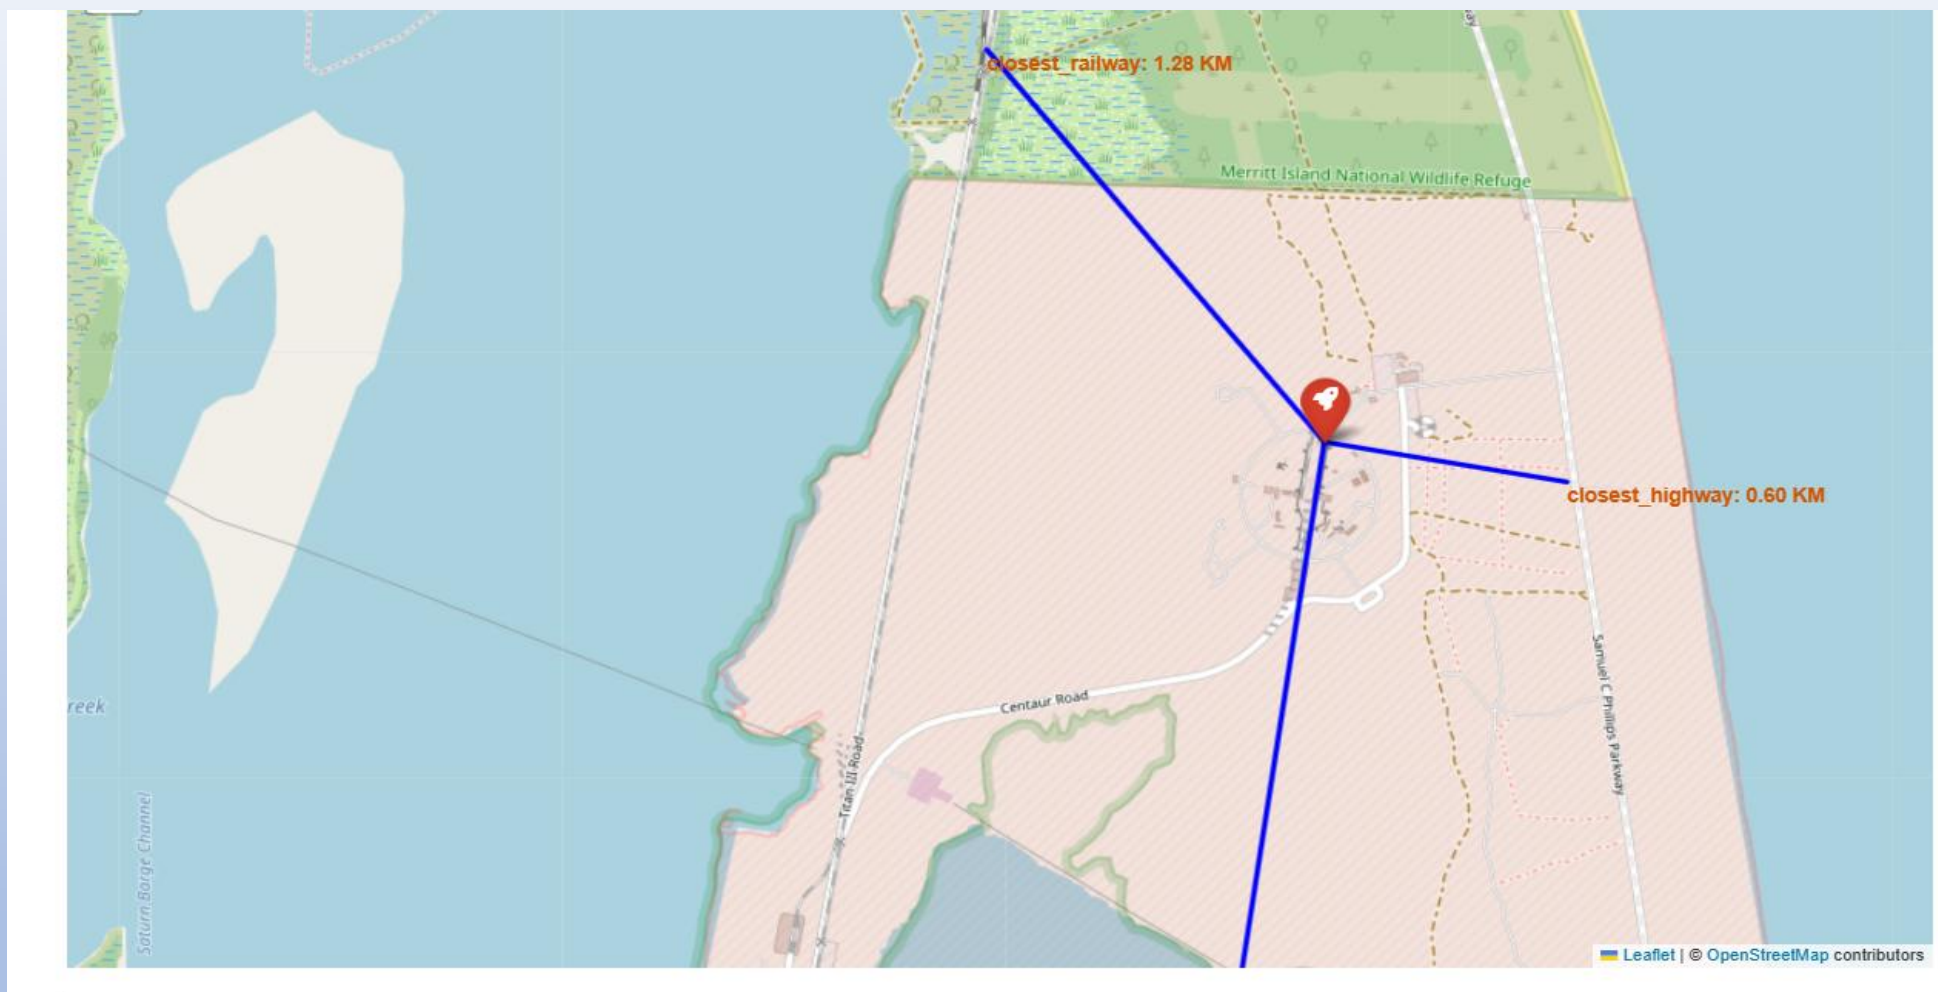













PLOTLY DASH-BOARD

SpaceX Launch Records Dashboard

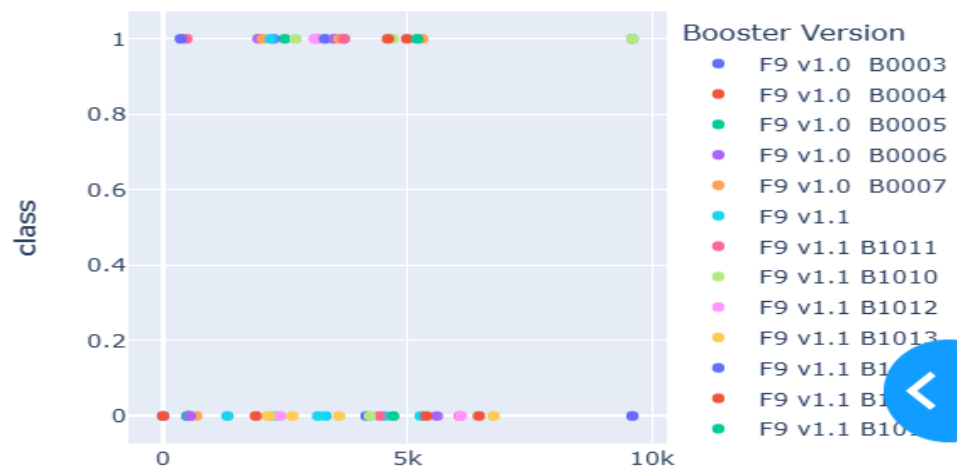
All Sites



Launch Successes and Failures for All Sites



Correlation between Payload and Success for All Sites




SpaceX Launch Records Dashboard


CCAFS LC-40



Launch Successes and Failures for CCAFS LC-40



Correlation between Payload and Success for CCAFS LC




SpaceX Launch Records Dashboard

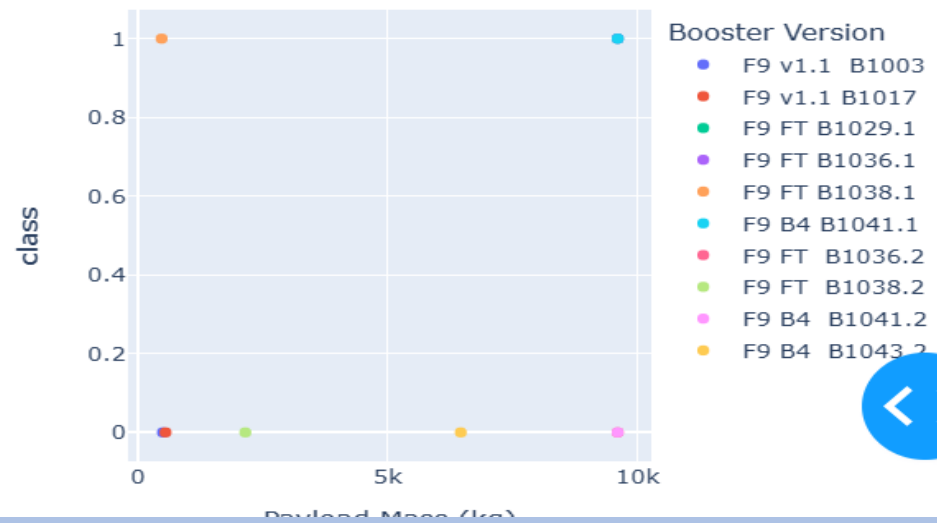
VAFB SLC-4E



Launch Successes and Failures for VAFB SLC-4E



Correlation between Payload and Success for VAFB SLC-4E




SpaceX Launch Records Dashboard


KSC LC-39A



Launch Successes and Failures for KSC LC-39A



Correlation between Payload and Success for KSC LC-39A




SpaceX Launch Records Dashboard

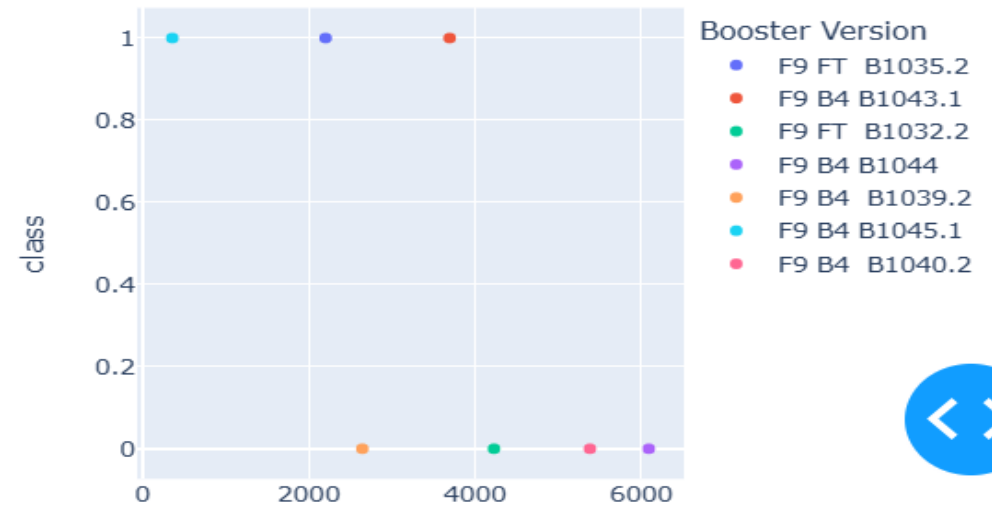
CCAFS SLC-40



Launch Successes and Failures for CCAFS SLC-40



Correlation between Payload and Success for CCAFS SL



PREDICTIVE ANALYSIS RESULTS

TASK 3

Use the function `train_test_split` to split the data `X` and `Y` into training and test data. Set the parameter `test_size` to 0.2 and `random_state` to 2. The training data and test data should be assigned to the following labels.

`X_train, X_test, Y_train, Y_test`

```
from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)

# Optionally, display the shapes of the outputs to confirm the split
print("Training data shape (X_train):", X_train.shape)
print("Test data shape (X_test):", X_test.shape)
print("Training labels shape (Y_train):", Y_train.shape)
print("Test labels shape (Y_test):", Y_test.shape)
```

```
Training data shape (X_train): (72, 84)
Test data shape (X_test): (18, 84)
Training labels shape (Y_train): (72,)
Test labels shape (Y_test): (18,)
```

we can see we only have 18 test samples.

TASK 4

Create a logistic regression object then create a GridSearchCV object `logreg_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

We output the `GridSearchCV` object for logistic regression. We display the best parameters using the data attribute `best_params_` and the accuracy on the validation data using the data attribute `best_score_`.

```
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

# Define the logistic regression model
lr = LogisticRegression()

# Parameters dictionary for GridSearchCV
parameters = {'C': [0.01, 0.1, 1],
              'penalty': ['l2'],
              'solver': ['lbfgs']}

# Create a GridSearchCV object with the logistic regression model
logreg_cv = GridSearchCV(lr, parameters, cv=10)

# Fit GridSearchCV to the training data
logreg_cv.fit(X_train, Y_train)

# Display the best hyperparameters and the accuracy of the best model
print("Tuned hyperparameters (best parameters): ", logreg_cv.best_params_)
print("Accuracy on validation data: ", logreg_cv.best_score_)
```

```
Tuned hyperparameters (best parameters): {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}
Accuracy on validation data: 0.8464285714285713
```

TASK 5



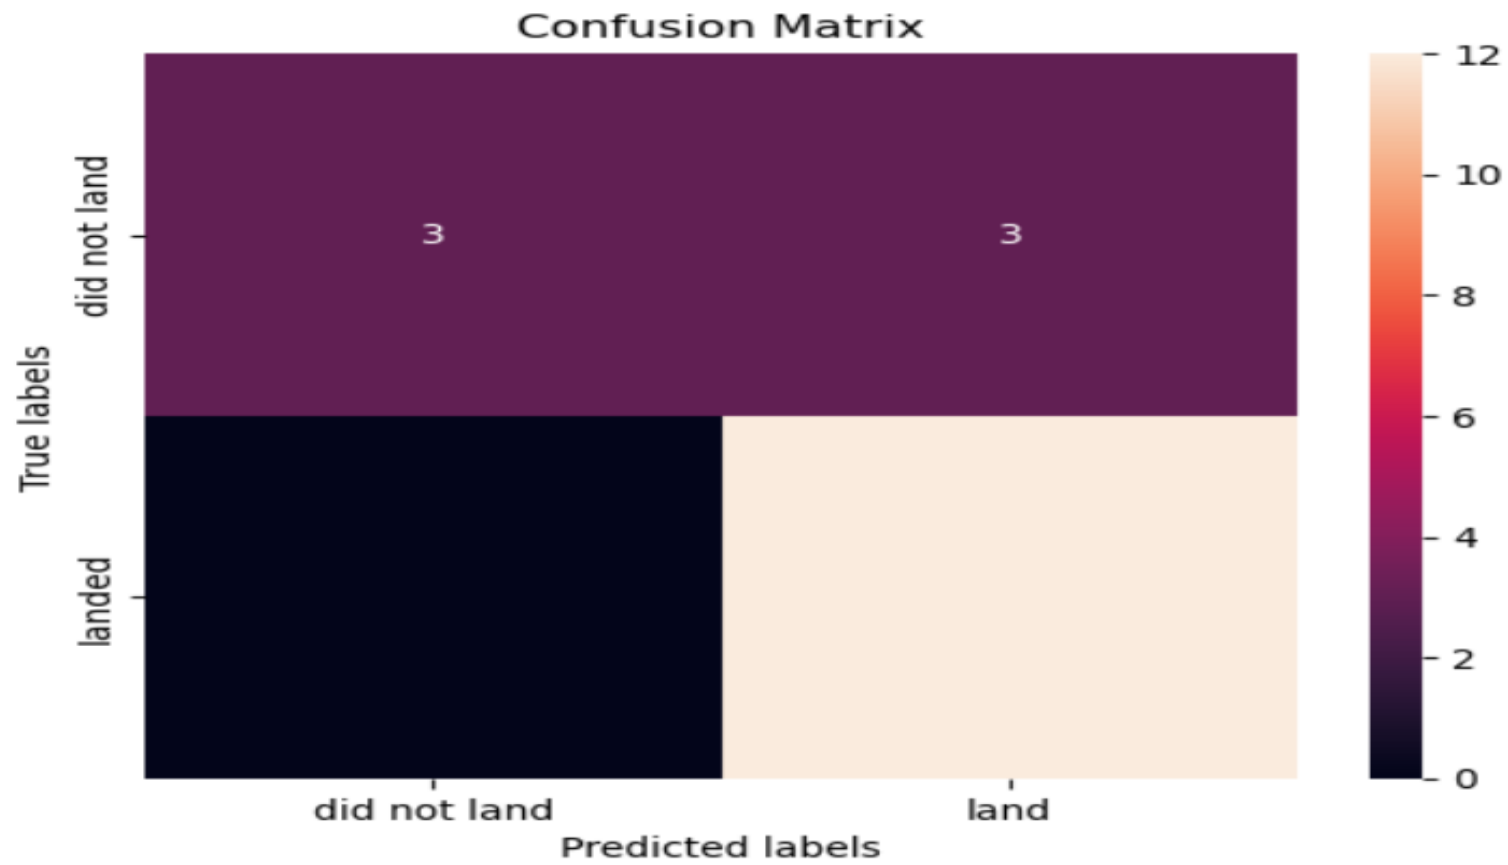
Calculate the accuracy on the test data using the method `score` :

```
: # Calculate the accuracy on the test data
test_accuracy = logreg_cv.score(X_test, Y_test)

# Print the accuracy on the test data
print("Accuracy on test data: ", test_accuracy)
```

Accuracy on test data: 0.8333333333333334


```
[24]: yhat=logreg_cv.predict(X_test)  
      plot_confusion_matrix(Y_test,yhat)
```



TASK 6

Create a support vector machine object then create a `GridSearchCV` object `svm_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters`.

```
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
import numpy as np

# Define the support vector machine model
svm = SVC()

# Parameters dictionary for GridSearchCV
parameters = {
    'kernel': ('linear', 'rbf', 'poly', 'sigmoid'),
    'C': np.logspace(-3, 3, 5),
    'gamma': np.logspace(-3, 3, 5)
}

# Create a GridSearchCV object with the support vector machine model
svm_cv = GridSearchCV(svm, parameters, cv=10)

# Fit GridSearchCV to the training data
svm_cv.fit(X_train, Y_train)

# Display the best hyperparameters and the accuracy of the best model
print("Tuned hyperparameters (best parameters): ", svm_cv.best_params_)
print("Accuracy on validation data: ", svm_cv.best_score_)
```

```
Tuned hyperparameters (best parameters): {'C': 0.03162277660168379, 'gamma': 0.001, 'kernel': 'linear'}
Accuracy on validation data: 0.8214285714285714
```

TASK 7

Calculate the accuracy on the test data using the method `score` :

```
# Calculate the accuracy on the test data  
test_accuracy = svm_cv.score(X_test, Y_test)  
  
# Print the accuracy on the test data  
print("Accuracy on test data: ", test_accuracy)
```

Accuracy on test data: 0.8333333333333334


```
# Expanding the parameter grid slightly more than the small test
expanded_parameters = {
    'criterion': ['gini', 'entropy'],
    'splitter': ['best', 'random'],
    'max_depth': [2, 4, 6],
    'max_features': ['auto', 'sqrt'],
    'min_samples_leaf': [1],
    'min_samples_split': [2, 5]
}

# Using more cross-validation folds than the smallest test, but fewer than the original large grid
expanded_tree_cv = GridSearchCV(tree, expanded_parameters, cv=5, verbose=3) # Moderate verbosity
expanded_tree_cv.fit(X_train, Y_train)
print("Expanded Grid Search Complete.")
print("Best parameters: ", expanded_tree_cv.best_params_)
print("Accuracy on validation data: ", expanded_tree_cv.best_score_)
```

#Expanded Grid Search Complete.

#Best parameters: {'criterion': 'gini', 'max_depth': 2, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 5, 'splitter': 'best'}

Accuracy on validation data: 0.8885714285714286



TASK 10

Create a k nearest neighbors object then create a `GridSearchCV` object `knn_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters` .

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV

# Define the K-Nearest Neighbors classifier
KNN = KNeighborsClassifier()

# Parameters dictionary for GridSearchCV
parameters = {
    'n_neighbors': list(range(1, 11)), # Using a list of neighbors from 1 to 10
    'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
    'p': [1, 2] # 1 corresponds to Manhattan distance, 2 to Euclidean distance
}

# Create a GridSearchCV object with the KNN model
knn_cv = GridSearchCV(KNN, parameters, cv=10)

# Fit GridSearchCV to the training data
knn_cv.fit(X_train, Y_train)

# Display the best hyperparameters and the accuracy of the best model
print("Tuned hyperparameters (best parameters): ", knn_cv.best_params_)
print("Accuracy on validation data: ", knn_cv.best_score_)
```

```
Tuned hyperparameters (best parameters): {'algorithm': 'auto', 'n_neighbors': 10, 'p': 1}
Accuracy on validation data: 0.8482142857142858
```

TASK 11

Calculate the accuracy of knn_cv on the test data using the method `score` :

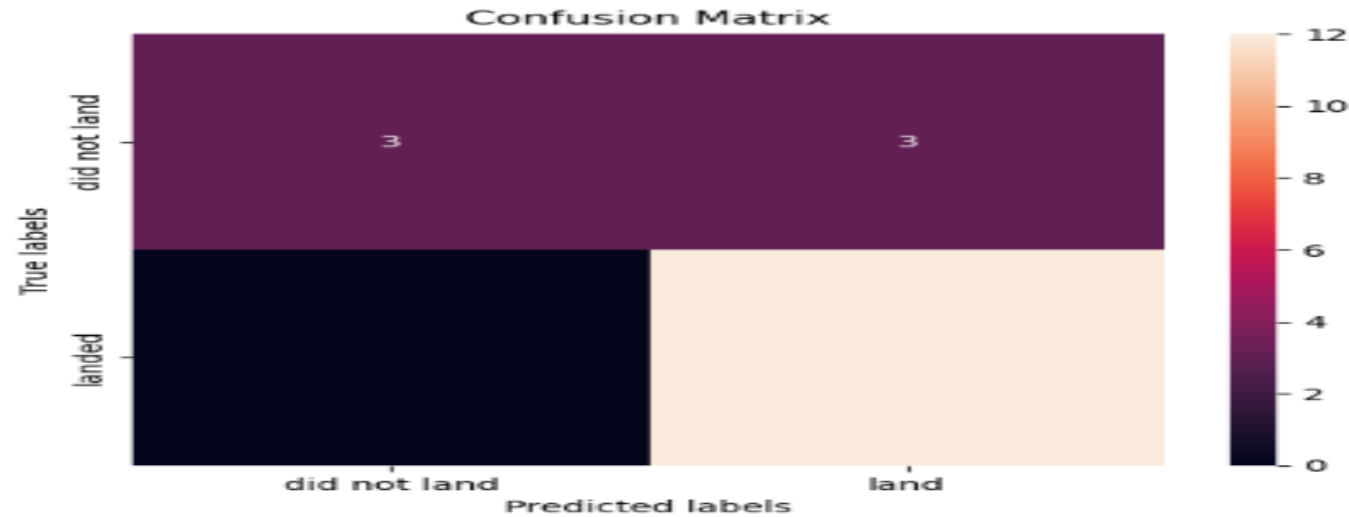
```
42]: # Calculate the accuracy on the test data
test_accuracy = knn_cv.score(X_test, Y_test)

# Print the accuracy on the test data
print("Accuracy on test data: ", test_accuracy)
```

Accuracy on test data: 0.8333333333333334

We can plot the confusion matrix

```
43]: yhat = knn_cv.predict(X_test)
plot_confusion_matrix(Y_test, yhat)
```



Conclusion – Exploratory Data Analysis(EDA)

Key Insights from Launch Data Analysis:

Launch Site Variation:

Success rates differ across launch sites, with some locations consistently achieving higher landing success, indicating site-specific operational advantages.

Payload Influence:

Heavier payloads, especially at certain launch sites, have been linked to higher landing success, suggesting a possible correlation between payload mass and performance.

Technological Progress:

A noticeable upward trend in success rates over time reflects improvements in rocket technology, engineering practices, and mission execution.

Orbit Type Success Rates:

Missions targeting certain orbit types tend to have higher success rates. This could be due to the relative complexity of the orbit or the organization's specialized expertise in specific mission types.

Flight Number Trends:

As the number of flights increases, so does the landing success rate. This pattern emphasizes the positive impact of accumulated experience and continuous process refinement.

Conclusion – Interactive Folium Maps

Strategic Placement of Launch Sites:

Coastal Proximity for Safety:

Launch sites are deliberately located near coastlines to ensure rockets have flight paths over open ocean. This minimizes the risk to populated areas in case of malfunctions or failures during launch.

Infrastructure Accessibility:

Mapping analysis reveals that launch facilities are positioned with easy access to key infrastructure such as roads and railways. This strategic location supports the efficient transportation of rocket parts and streamlines logistical operations crucial for timely and secure launch preparations.

Conclusion – Plotly Dash-Boards

Launch Outcomes by Site:

Analysis reveals differences in success rates across various launch sites. Some locations consistently report higher success, indicating potential site-specific advantages or more optimized operational procedures, while others may face unique challenges affecting outcomes.

Payload and Success Correlation:

Scatter plot analysis indicates no strong correlation between payload mass and launch success. This suggests that SpaceX's launch systems perform reliably across a broad spectrum of payload weights, highlighting the robustness and adaptability of their rocket technology.

Conclusion – Predictive Analysis

Model Accuracy:

All evaluated models achieved accuracy scores exceeding 80%, demonstrating strong predictive performance on test data. This indicates that the selected features play a significant role in determining the outcome of Falcon 9 landings.

Consistency Across Models:

Algorithms such as logistic regression, support vector machines, and k-nearest neighbors produced comparable results, suggesting that the dataset is well-structured and contains clear patterns that multiple models can effectively learn.

Confusion Matrix Insights:

The confusion matrices reveal high counts of true positives and true negatives, indicating reliable model predictions for both successful and failed landings. However, the occurrence of false positives and false negatives highlights areas where the models can still be refined.

Hyperparameter Optimization with GridSearchCV:

Tuning model parameters using GridSearchCV led to improved accuracy scores compared to using default settings. This underscores the importance of hyperparameter optimization in enhancing model performance and reliability.

Creativity Applied in Presentation

Effective Presentation Design:

Designing presentations with dynamic layouts and visually appealing backgrounds transforms raw data into an engaging visual story.

Clarity Through Simplicity:

Use concise bullet points and structured content to make complex ideas easy to follow.

Visual Impact:

Creative slide elements not only capture attention but also improve understanding and memory retention.

This approach bridges the gap between data and audience, making your message both impactful and memorable.

Innovative Insights from the Data Analysis

1. Insight from Exploratory Data Analysis (EDA):

A compelling observation from the visualizations is the steady improvement in launch success rates over time, reflecting a clear learning curve. As the number of SpaceX flights increases, so does the rate of successful landings. This trend highlights the iterative nature of aerospace engineering—where each mission generates critical feedback, driving continuous technological refinement and operational efficiency in future launches.

2. Insight from Predictive Analysis:

A key takeaway from the predictive modeling phase is the substantial impact of hyperparameter tuning on model performance. Regardless of the algorithm used, the application of GridSearchCV consistently optimized parameters and improved accuracy. This underscores the power of automated machine learning (AutoML) techniques in enhancing model outcomes, reducing the need for manual adjustments, and promoting scalable, data-driven solutions.

Thank you